

Dynamic Compilation of DSLs

Vojin Jovanovic

EPFL, Switzerland
vojin.jovanovic@epfl.ch

Domain-specific language (DSL) compilers use *domain knowledge* to perform *domain-specific optimizations* that can yield several orders of magnitude speedups [4]. These optimizations, however, often require knowledge of the run-time values. For example, in matrix chain multiplication, knowing sizes of matrices allows choosing the optimal multiplication order [2, Ch. 15.2] and in relational algebra knowing data sizes is necessary for choosing the right join order [5]. As a motivating example we will use multiplication of three matrices in a DSL for linear algebra embedded in Scala:

```
val (m1, m2, m3) = ... // matrices of unknown size
m1 * m2 * m3
```

In this program, without knowing the matrix sizes, the DSL compiler can not determine the optimal order of multiplications. There are two possible orders $(m1*m2)*m3$ with an estimated cost $c1$ and $m1*(m2*m3)$ with an estimated cost $c2$ where:

```
c1 = m1.rows*m1.columns*m2.columns+m1.rows*m2.columns*m3.rows
c2 = m2.rows*m2.columns*m3.columns+m1.rows*m2.rows*m3.columns
```

Ideally we would change the multiplication order at runtime only when the condition $c1 > c2$ changes. The only solution for producing an optimal program is to use *dynamic compilation* [1].

Existing dynamic compilation systems, however, are not a right fit for this task as they use run-time information for transformations that are not domain specific, e.g., specialization [3], loop unrolling [6], etc. As a consequence:

- Value profiling tracks the *stability* of values. In our example we need to track stability over computation over several (possibly unstable) values that are used for computing $c1$ and $c2$.
- Guards for recompilation are implemented as equality checks between the current value and the previous stable value of the variable. In our example we need comparisons with previous versions of the computation over run-time values ($c1 > c2$).
- Code caches are based on checking equality of stable values. In our example, we need equality with computed values that change the outcome of DSL compilation.

Only, exception to existing dynamic compilation systems is Truffle [7] that allows creation of custom profiling and recompilation guards. However, with Truffle, language designers do not have the full view of the program, and thus,

can not perform global optimizations (e.g., matrix chain multiplication optimization).

We propose a dynamic compilation system aimed for domain specific languages in which:

- DSL authors declaratively state the values that are of interest for profiling (e.g., array and matrix sizes, vector and matrix sparsity, etc.).
- The DSL compiler is instrumented to reify all operations and transformation decisions based on the profiled run-time values. In our example, all computations on run-time values performed during matrix chain multiplication optimization would be reified and stored (i.e., `c1` and `c2`).
- Automatically introduces guards based on the stored compilation process (i.e., `c1 > c2` in the example).
- Automatically adds code caches that are not addressed by stable values but with outcomes of the DSL compilation decisions. In the example the code cache would have two entries addressed with a single boolean value `c1>c2`.

In the talk we will present the prototype of the dynamic compilation system for domain specific languages and the performance results on several examples in linear algebra. Our goal is to foster discussion on different policies for the automatic introduction of recompilation guards: *i)* heuristic, *ii)* *DSL author specified*, and *cost based*.

References

1. Joel Auslander, Matthai Philipose, Craig Chambers, Susan J Eggers, and Brian N Bershad. Fast, effective dynamic compilation. In *International Conference on Programming Language Design and Implementation (PLDI)*, 1996.
2. Thomas H Cormen, Charles E Leiserson, Ronald L Rivest, Clifford Stein, et al. *Introduction to algorithms*, volume 2. MIT press Cambridge, 2001.
3. Brian Grant, Markus Mock, Matthai Philipose, Craig Chambers, and Susan J Eggers. DyC: an expressive annotation-directed dynamic compiler for C. *Theoretical Computer Science*, 248(1), 2000.
4. Tiark Rompf, Arvind K. Sujeeth, Nada Amin, Kevin J. Brown, Vojin Jovanović, HyoukJoong Lee, Manohar Jonnalagedda, Kunle Olukotun, and Martin Odersky. Optimizing data structures in high-level programs: New directions for extensible compilers based on staging. In *Symposium on Principles of Programming Languages (POPL)*, 2013.
5. P Griffiths Selinger, Morton M Astrahan, Donald D Chamberlin, Raymond A Lorie, and Thomas G Price. Access path selection in a relational database management system. In *International conference on Management of data (SIGMOD)*, pages 23–34, 1979.
6. Toshio Sukanuma, Toshiaki Yasue, Motohiro Kawahito, Hideaki Komatsu, and Toshio Nakatani. A dynamic optimization framework for a java just-in-time compiler. In *International Conference on Programming Language Design and Implementation (PLDI)*, volume 36, pages 180–195. ACM, 2001.

7. Thomas Würthinger, Christian Wimmer, Andreas Wöß, Lukas Stadler, Gilles Duboscq, Christian Humer, Gregor Richards, Doug Simon, and Mario Wolczko. One VM to rule them all. In *Symposium on New Ideas, New Paradigms, and Reflections on Programming & Software (Onward!)*, 2013.