

Type-Driven Partial Evaluation without Code Duplication

No Author Given

No Institute Given

Abstract.

Keywords: Partial Evaluation

1 Introduction

Partial evaluation [5] is an optimization technique that identifies *statically known* program parts and pre-computes them at compile time. Partial evaluation has been intensively studied and successfully applied: for removing abstraction overheads in high-level programs [2, 8], for domain-specific languages [1, 6], and for converting language interpreters into compilers [4, 9, 12]. Applying partial evaluation in these domains often improves program performance by several orders of magnitude [10, 1].

To achieve *predictable* and *safe* partial evaluation, however, a partial evaluator must be instructed by the programmer [1, 7]. Unlike other compiler optimizations, due to compile-time execution, partial evaluation might not *terminate*. Furthermore, *code explosion* is possible as the final program can be arbitrarily large due to compile-time execution. Lack of predictability and danger of code explosion are the reason that successful partial evaluators [1, 11, 8, 12] are programmer controlled.

To show how programmers can control partial evaluation we define a function `dot` for computing a dot-product of two vectors that contain numeric values.

```
def dot[V:Numeric](v1: Vector[V], v2: Vector[V]): V =  
  (v1 zip v2).foldLeft(zero[V]) { case (prod, (c1, cr)) =>  
    prod + c1 * cr  
  }
```

In order to remove the abstraction overhead of `zip` and `foldLeft` the partial evaluator must apply extensive analysis to conclude that vectors are static in size and that this can be later used to unroll the recursion inside `foldLeft`. Even if the analysis is successful the evaluator must be conservative about unrolling the `foldLeft` as vector sizes, and thus the produced code, can be very large. What

⁰ We use Scala for all code examples in this paper. In order to comprehend the paper the reader is required to know the mere basics of the language

if we know that vector sizes are relatively small and we would like to predictably unroll `dot` into a flat sum of products?

Ideally, the programmer would with a minimal number of annotations be able to: *i)* require that `v1 v2` vectors are of statically known size, *ii)* require that all operations on vector arguments should be further partially evaluated, *iii)* allow elements of vectors to be generic, and *iv)* reuse the existing implementation of the `Vector` data structure.

[TODO: expand] The main idea of this paper is to explicitly capture the user intent about partial evaluation in the types. We annotate every type in the language with one of the three values:

- *dynamic* signifies that the value of the type is not known at compile-time. In code *dynamic* is represented as `@d`.
- *static* signifies that the value is known at compile-time. In code *static* is represented as `@s`.
- *inline* requires that the type is statically known and guarantees that operations on the term will be partially evaluated. In code *inline* is represented as `@i`.

With our partial evaluator, to require that vectors `v1` and `v2` are static and to partially evaluate the function, a programmer would need to make a simple modification of the `dot` signature:

```
def dot[V: Numeric](v1: Vector[V] @i, v2: Vector[V] @i): V
```

This, in effect, requires that only vector arguments (not their elements) are statically known and that all operations will be inlined and further partially evaluated. Residual programs of `dot` application in different cases are:

```
// [el1, el2, el3, el4] are dynamic
dot(Vector(el1, el2), Vector(el3, el4))
  ⇨ el1 * el3 + el2 * el4

dot(Vector(2, 4), Vector(1, 10))
  ⇨ 2 * 1 + 4 * 10

// inline promotes static terms into inline
dot(Vector(inline(2), inline(4)), Vector(inline(1), inline(10)))
  ⇨ 42
```

Predictable partial evaluation has been a goal of many projects (§6) which can be roughly categorized into following categories:

- Programming languages Idris and D provide annotations `static` that can be placed on function arguments. In these systems static parameters are required to be deeply static and, thus, the `dot` function could not be expressed.
- Type-directed partial evaluation [3] and LMS [8] use types to communicate the programmers intent about partial evaluation and, thus, can express `dot` function. These approaches, however, require the programmer to implement two versions of `Vector` and other operations. This fosters, costly and hardly maintainable, code duplication.

– MetaOCaml [11]

Contributions:

Evaluation:

Sections:

2 Formalization $F_{i<}$:

$t ::=$	Terms:
x, y	identifier
$(x : iT) \Rightarrow t$	function
$t(t)$	application
$\{x = t\}$	record
$t.x$	selection
$inline\ t$	inlining request
$[X <: iT] \Rightarrow t$	type abstraction
$t[iT]$	type application
$S, T, U ::=$	Types:
$iS \Rightarrow jT$	function type
$\{x : iS\}$	record type
$[X <: iS] \Rightarrow jT$	universal type
\top	top type
$iT, jT, kT, lT ::=$	Inlineable Types:
$T, dynamic\ T$	dynamic type
$static\ T$	static type
$inline\ T$	must inline type
$\Gamma ::=$	Contexts:
\emptyset	empty context
$\Gamma, x : iT$	term binding
$\Gamma, X <: iT$	type binding

Fig. 1. Syntax

$\frac{x : iT \in \Gamma}{\Gamma \vdash x : iT}$	(T-IDENT)
$\frac{\Gamma, x : iT_1 \vdash t : jT_2}{\Gamma \vdash (x : iT_1) \Rightarrow t : \text{static } iT_1 \Rightarrow jT_2}$	(T-FUNC)
$\frac{\Gamma \vdash t : iT}{\Gamma \vdash \{x = t\} : \text{static } \{x : iT\}}$	(T-REC)
$\frac{\Gamma \vdash t_1 : i(jT_1 \Rightarrow kT_2) \quad \Gamma \vdash t_2 : jT_2}{\Gamma \vdash t_1(t_2) : (i \wedge j \wedge k)T_2}$	(T-APP)
$\frac{\Gamma \vdash t : i\{x = jT_1, y = kT_2\}}{\Gamma \vdash t.x : (i \wedge j)T_1}$	(T-SEL)
$\frac{\Gamma \vdash t : \text{static } T}{\Gamma \vdash \text{inline } t : \text{inline } T}$	(T-INLINE)
$\frac{\Gamma, X <: iT_1 \vdash t_2 : jT_2}{\Gamma \vdash [X <: iT_1] \Rightarrow t_2 : \text{static } [X <: iT_1] \Rightarrow jT_2}$	(T-TABS)
$\frac{\Gamma \vdash t_1 : i([X <: jT_{11}] \Rightarrow kT_{12}) \quad \Gamma \vdash lT_2 <: jT_{11}}{\Gamma \vdash t_1[lT_2] : [X \mapsto lT_2](i \wedge k)T_{12}}$	(T-TAPP)
$\frac{\Gamma \vdash t : iS \quad \Gamma \vdash iS <: jT}{\Gamma \vdash t : jT}$	(T-SUB)

Fig. 2. typing $\Gamma \vdash t : iT$

$\frac{\forall i. i <: \text{dynamic}}{\forall i \in \{\text{static}, \text{inline}\}. i <: \text{static}}$	(IS-DYNAMIC)
$\frac{\text{inline } <: \text{inline}}{\text{inline } <: \text{inline}}$	(IS-STATIC)
	(IS-INLINE)

Fig. 3. Inlinity Subtyping $i <: j$

$$\frac{\forall i, \forall j, i <: j. i \wedge j = j}{\forall i, \forall j, i <: j. i \wedge j = j}$$

Fig. 4. Inlinity Intersection $i \wedge j$

$$\begin{array}{c}
\frac{}{\Gamma \vdash S <: S} \quad (\text{S-REFL}) \\
\frac{\Gamma \vdash S <: U \quad \Gamma \vdash U <: T}{\Gamma \vdash S <: T} \quad (\text{S-TRANS}) \\
\frac{}{\Gamma \vdash S <: \top} \quad (\text{S-TOP}) \\
\frac{X <: T \in \Gamma}{\Gamma \vdash X <: T} \quad (\text{S-TVAR}) \\
\frac{\Gamma \vdash kT_1 <: iS_1 \quad \Gamma \vdash jS_2 <: lT_2}{\Gamma \vdash iS_1 \Rightarrow jS_2 <: kT_1 \Rightarrow lT_2} \quad (\text{S-ARROW}) \\
\frac{\Gamma, X <: iU_1 \vdash jS_2 <: kT_2}{\Gamma \vdash [X <: iU_1] \Rightarrow jS_2 <: [X <: iU_1] \Rightarrow kT_2} \quad (\text{S-ALL}) \\
\frac{\{x_p : i_p T_p^{p \in 1..n+m}\} <: \{x_p : i_p T_p^{p \in 1..n}\}}{\forall p \in 1..n. i_p S_p <: j_p T_p} \quad (\text{S-RECDWIDTH}) \\
\frac{\{x_p : i_p S_p^{p \in 1..n}\} <: \{x_p : j_p T_p^{p \in 1..n}\}}{\{x_p : i_p S_p^{p \in 1..n}\} \text{ is permutation of } \{y_p : j_p T_p^{p \in 1..n}\}} \quad (\text{S-RECDDEPTH}) \\
\frac{\{x_p : i_p S_p^{p \in 1..n}\} \text{ is permutation of } \{y_p : j_p T_p^{p \in 1..n}\}}{\{x_p : i_p S_p^{p \in 1..n}\} <: \{y_p : j_p T_p^{p \in 1..n}\}} \quad (\text{S-RECPERM}) \\
\frac{i <: j \quad \Gamma \vdash S <: T}{\Gamma \vdash iS <: jT} \quad (\text{S-INLINE})
\end{array}$$

Fig. 5. Subtyping $\Gamma \vdash iS <: jT$

3 Translating Scala to the Core Calculus

4 Case Studies

4.1 Integer Power Function

- Explain what happens.
- Typical partial evaluation example. Can be handled by D and Idris and not without duplication with type-driven partial evaluation.

```

@i? def pow(base: Double, exp: Int @i?): Double =
  if (exp == 0) 1 else base * pow(base, exp)

```

Fig. 6. Function for computing the non-negative power of a real number.

4.2 Variable Argument Functions

- *@i?* in argument position is a macro that expands the function to an underlying function *@i?* def min_underlying[T: Numeric](values: Seq[T] *@i?*): T and a macro that will call it according to the input parameters.
- Comparison to other approaches.

```
@i? def min[T: Numeric](@i? values:T*): T =
  values.tail.foldLeft(values.head)((min, el) => if (el < min) el else min)
```

Fig. 7. Function for computing the non-negative power of a real number.

4.3 Butterfly Networks

- Reference LMS. Discuss @i! annotation on classes. Works for both dynamic and static inputs.
- Comparison to LMS. Mention a pervasive number of annotations. Discuss duality of Exp[T] and @i!.

4.4 Dot Product

- Explain the removal of type classes together with inline. Explain how type classes are @i? and how they will completely evaluate if they are passed a static value.
- Comparison to other approaches.

```
object Numeric {
  @i! implicit def dnum: Numeric[Double] @i! = DoubleNumeric
  @i! def zero[T](implicit num: Numeric[T]): T = num.zero
  object Implicit {
    @i! implicit def infixNumericOps[T](x: T)(implicit num: Numeric[T]): Numeric[T]#Ops = new
  }
}

trait Numeric[T] {
  def plus(x: T, y: T): T
  def times(x: T, y: T): T
  def zero: T

  @i! class Ops(lhs: T) {
    @i! def +(rhs: T) = plus(lhs, rhs)
    @i! def *(rhs: T) = times(lhs, rhs)
  }
}

object DoubleNumeric extends Numeric[Double] {
  @i! def plus(x: Double @i?, y: Double @i?): Double = x + y
  @i! def times(x: Double @i?, y: Double @i?): Double = x * y
  @i! def zero: Double = 0.0
}
```

Fig. 8. Function for computing the non-negative power of a real number.

5 Evaluation

6 Related Work

7 Conclusion

References

1. Edwin C. Brady and Kevin Hammond. Scrapping your inefficient engine: Using partial evaluation to improve domain-specific language implementation. In *International Conference on Functional Programming (ICFP)*, 2010.
2. Jacques Carette and Oleg Kiselyov. Multi-stage programming with functors and monads: Eliminating abstraction overhead from generic code. In *Generative Programming and Component Engineering (GPCE)*, 2005.
3. Olivier Danvy. *Type-directed partial evaluation*. Springer, 1999.
4. Yoshihiko Futamura. Partial evaluation of computation process—an approach to a compiler-compiler. *Higher-Order and Symbolic Computation*, 12(4):381–391, 1999.
5. Neil D. Jones, Carsten K. Gomard, and Peter Sestoft. *Partial Evaluation and Automatic Program Generation*. Prentice Hall, 1993.
6. Manohar Jonnalagedda, Thierry Coppey, Sandro Stucki, Tiark Rompf, and Martin Odersky. Staged parser combinators for efficient data processing. In *International Conference on Object Oriented Programming Systems Languages and Applications (OOPSLA)*, 2014.
7. Anne-Françoise Le Meur, Julia L Lawall, and Charles Consel. Specialization scenarios: A pragmatic approach to declaring program specialization. *Higher-Order and Symbolic Computation*, 17(1-2):47–92, 2004.
8. Tiark Rompf and Martin Odersky. Lightweight modular staging: a pragmatic approach to runtime code generation and compiled DSLs. *Communications of the ACM*, 55(6):121–130, June 2012.
9. Tiark Rompf, Arvind K. Sujeeth, Kevin J. Brown, HyoukJoong Lee, Hassan Chafi, Kunle Olukotun, and Martin Odersky. Project Lancet: Surgical precision JIT compilers. In *International Conference on Programming Language Design and Implementation (PLDI)*, 2013.
10. Amin Shali and William R. Cook. Hybrid partial evaluation. In *International Conference on Object Oriented Programming Systems Languages and Applications (OOPSLA)*, 2011.
11. Walid Taha and Tim Sheard. Multi-stage programming with explicit annotations. In *Workshop on Partial Evaluation and Program Manipulation (PEPM)*, 1997.
12. Thomas Würthinger, Christian Wimmer, Andreas Wöß, Lukas Stadler, Gilles Duboscq, Christian Humer, Gregor Richards, Doug Simon, and Mario Wolczko. One vm to rule them all. In *Symposium on New Ideas, New Paradigms, and Reflections on Programming & Software (Onward!)*, 2013.