Compile-Time Views: Predictable Type-Directed Partial Evaluation Without Code Duplication

Abstract

Categories and Subject Descriptors D.3.3 [*Programming Languages*]: Language Constructs and Features

Keywords Multi-Stage Programming, Partial Evaluation

1. Introduction

Multi-stage programming (or *staging*) is a flavor of meta-programming where compilation is separated in multiple *stages*. Execution of each stage outputs code that is executed in the next stage of compilation. The first stage of compilation is the *host language* compile-time, the second stage the host language runtime, the third stage is the run-time of run-time generated code, etc. Notable, staging frameworks are MetaOCaml [13] and LMS [10], and they were successfully applied as a *partial evaluatior* [4]: for removing abstraction overheads in high-level programs [1, 10], for domain-specific languages [5], and for converting language interpreters into compilers [3, 11].

We show an example of how staging is used for partial evaluation of a function for computing the inner product of two vectors [1]:

```
def dot[V:Numeric](v1: Vector[V], v2: Vector[V]): V =
  (v1 zip v2).foldLeft(zero[V]) {
    case (prod, (cl, cr)) => prod + cl * cr
  }
```

In function `dot`, if the vector sizes are static during program runtime the inner product can be partially evaluated into a sum of products of vector components. To achieve partial evaluation, we must communicate to the staging framework that the operations on values of vector components

should be executed in the next stage (after language runtime). The compilation stage in which a term will be executed is determined by *code quotation* (in MetaOCaml) or by specific parametric types `Rep` (in LMS). In LMS [2] we mark that the vector size is statically known by annotating only vector elements with a `Rep` type:

```
def dot[V:Numeric]
  (v1: Vector[Rep[V]], v2: Vector[Rep[V]]): Rep[V]
```

Here the `Rep` annotations on `V` denote that elements of vectors will be known only in the next stage (after run-time compilation). At this stage the `zip foldLeft`, and pattern matching inside the closure will not exist as they were evaluated at the previous stage (host language runtime). Note that the unquoted/unannotated code is always executed during host-language runtime and quoted/type annotated code is executed after run-time compilation.

First Problem. How can we use staging for programs whose values are statically known at the host language compile-time (the first stage)? All staging frameworks treat unannotated terms as host language runtime values and annotated terms as values of later stages. Even if we would start staging one stage earlier (at host language compile-time), we would have to annotate all run-time values. Annotating all values is cumbersome since host language run-time values comprise the majority of user programs **[TODO: cf.]**.

Programming languages Idris and D allow try to solve this problem by allowing the `static` annotation on function arguments. Annotation `static` denotes that the term is statically known and that all operations on that term should be executed at compile-time. However, since `static` is placed on terms rather then types, it can mark only *whole terms* as static. This restricts the number of programs that can be expressed, *e.g.*, we could not express that vectors in the signature of *dot* are static only in size. Finally, information about `static` terms can not be propagated through return types of

---

[1] All code examples are written in *Scala*. For comprehension of the paper basic knowledge of Scala is necessary.

[2] In this work we use LMS as it is the only staging framework in Scala.

functions, so `static` in Idris and D is more a partial evaluation construct.

Second Problem. Staging systems based on type annotations (e.g., LMS and type-directed partial evaluation [2]) inherently require duplication of code as ,a priory, no operations are defined on `Rep` annotated types. For example, in the `dot` function all numerical types (*e.g.*, `Rep[Int]`, `Rep[Double]`, etc.) must be re-implemented in order to typecheck the programs and achieve code generation for the next stage.

Suereth et al. [12] and Jovanovic et al. [6] propose generating code for the next stage computations based on a language specification. These approaches solve the problem, but they require writing additional specification for the libraries, require a large machinery for code generation, and support only restricted parts of Scala.

The main idea of this paper is that annotated types should denote computations that happen at the *previous stage* instead of the next stage. The reason is two-fold: *i)* annotating code of previous stages succinctly express compile-time execution and *ii)* in staged programs the static terms appear less frequently than run-time terms, and in order to bear minimum overhead for the users, it is better to add annotation overhead to static terms.

Further, annotated types are simply a*compile-time view* of existing data types and therefore no code duplication is necessary. The compile-time view makes all operations and non-generic fields executed in the host language compile time. The compile-time view requires programmers to define a single definition of a type. Then, the existing types can be promoted to their compile-time duals with the `@ct` annotation at the type level, and with the `ct` function on the term level. Consequently, due to the integration with the type system, the control over staging is fine-grained and polymorphic, and term level promotions obviate code duplication for static data structures.

With compile-time views, to require that vectors `v1` and `v2` are static and to partially evaluate the function, a programmer would need to make a simple modification of the `dot` signature:

```
def dot[V: Numeric]
  (v1: Vector[V]@ct, v2: Vector[V]@ct): V
```

This, in effect, requires that only vector arguments (not their elements) are statically known and that all operations on vector arguments will be executed at compile time (partially evaluated). Since, values are polymorphic the result of the function will either be a dynamic value, static value, or a compile-time value that can be further used for partial evaluation. Residual programs of calling `dot` with arguments from different stages:

```
// [el1, el2, el3, el4] are dynamic decimals
dot(ct(Vector)(el1, el2), ct(Vector)(el3, el4))
   ↪ (el1 * el3 + el2 * el4): Double@dynamic
```

```
dot(ct(Vector)(2.0, 4.0), ct(Vector)(1.0, 10.0))
   ↪ (2.0 * 1.0 + 4.0 * 10.0): Double@static
```

```
// ct promotes static terms to compile-time
dot(ct(Vector)(ct(2), ct(4)),
    ct(Vector)(ct(1), ct(10)))
   ↪ 42: Double@ct
```

In this paper we contribute to the state of the art:

- By introducing compile-time views as means to: *i)* succinctly achieve staging at host language compile-time and to *ii)* avoid code duplication in type based staging systems.

- By introducing the $F_{i<:}$ calculus (§5) that in a fine-grained way captures the user's intent about partial evaluation. The calculus is based on $F_{<:}$ with lazy records which makes it suitable for representing modern multi-paradigm languages with object oriented features. Finally, we formally define evaluation rules for $F_{i<:}$.

- By providing a *translation scheme* from data types in object oriented languages (polymorphic classes and methods) into their dual compile-time views in the $F_{i<:}$ calculus (§6).

- By demonstrating the usefulness of compile-time views in four case studies (§3): inlining, partially evaluating recursion, removing overheads of variable argument functions, and removing overheads of type-classes [9].

We have implemented ScalaCT according to the translation scheme (§6) from object oriented features of Scala to the $F_{i<:}$ calculus. The prototype implemented for Scala and open-sourced[3]. It has a minimal Scala interface (§2) based on type annotations. We have evaluated performance gains and the validity of the partial evaluator on all case studies (§3) and compared them to LMS. In all benchmarks our evaluator gives significant performance gains compared to original programs and performs equivalently to LMS.

2. Compile-Time Views for Scala

We have implemented ScalaCT, a staging extension for Scala based on compile-time views. ScalaCT is a compiler plugin that executes in a phase after the Scala type checker. The plugin starts with pre-typed Scala programs and uses type annotations [8] to track and verify information about the biding-time of terms. Currently, it supports only two stages of compilation: host language compile-time (types annotated with `@ct`) and host language run-time (unannotated code).

To the user, ScalaCT exposes a minimal interface (Figure 2) with annotations `inline` and `ct`, and functions `inline` and `ct`.

Annotation `ct` is used at the type level (*e.g.*, `Int@ct`) and denotes a compile-time view of a type. The annotation is integrated in the Scala's type system and, therefore, can be

---

[3] Source code: `https://github.com/scala-inline/scala-inline`.

```
package object scalact {

  final class ct extends StaticAnnotation
  final class inline extends StaticAnnotation

  @compileTimeOnly def ct[T](body: => T): T = ???
  @compileTimeOnly def inline[T](body: => T): T = ???

}
```

Figure 1. Interface of the ScalaCT.

arbitrarily nested in different variants of types. Table 2 shows how the `@ct` annotation can be placed on types and how it, due to the translation to the compile-time views (Figure **??**), changes method signatures on annotated types.

In Table 2, `Int@ct` is a non-polymorphic type and therefore according to the translation to the compile-time view (13) parameters of all methods will also be compile-time views. On the other hand, `Vector[Int]@ct` will have parameters of all methods transformed except the generic ones. In effect, this, makes higher order combinators of `Vector` operate on dynamic values, thus, function `f` passed to `map` accepts the dynamic value as input. Type `Vector[Int@ct]@ct` has all parts executed at compile-time. The return type of the function `map` can still be both dynamic and a compile-time view: due to the type parameter U.

Annotation inline can be used only at the term level on statically known methods and functions. It denotes that the method/function will be inlined during compilation time. In other words, `inline` is marking that the function application is a compile-time computation and that application should be removed by partial evaluation. This is not the first time that inlining is achieved through partial evaluation [7].

Internally `inline` can be expressed in terms of the `ct` annotation. A method

```
@inline def dot[V: Numeric]
  (v1: Vector[V], v2: Vector[V]): V
```

will have an internal method type

```
((v1: Vector[V], v2: Vector[V]) => V)@ct
```

that can not be written by the users. We choose the name `inline` to be consistent with the existing Scala `inline` annotation.

Functions `ct` **and** `inline` are used at the term level for promoting Scala objects and methods/functions into their compile-time views. Without the the `ct` and `inline` we would not be able to instantiate compile-time views of types. Table 2 shows how different types of terms are promoted to their compile-time views.

Function `ct` can be applied to objects (*e.g.*, `Vector`) to provide a compile-time view over their methods. When those objects have generic parameters, `ct` be used to promote the arguments, and thus, the result types of these functions.

When applied, on functions `ct` promotes the compile-time view as well as its arguments and the return type.

Function `inline` can be applied only on functions to promote only the function (and not the arguments) to their compile time views. This function can be seen as a shallow version of `ct` that makes only the outer type a compile-time view.

3. Case Studies

In this section we present selected use-cases for compile-time views that demonstrate the core functionality. We start with a canonical example of the power function (§3.2), then we demonstrate how variable argument functions can be desugared into the core functionality (§3.3). Finally, we demonstrate how the abstraction overhead of the `dot` function and all associated type-classes can be removed (§3.5).

3.1 Inlining Expressed Through Partial-Evaluation

3.2 Recursion

The canonical example in partial evaluation is the computation of the integer power function:

```
def pow(base: Double, exp: Int): Double =
  if (exp == 0) 1 else base * pow(base, exp)
```

When the exponent (`exp`) is statically known this function can be partially evaluated into `exp` multiplications of the `base` argument, significantly improving performance [].

With compile-time views making `pow` partially evaluated requires adding two annotations:

```
@inline def pow(base: Double, exp: Int @ct): Double =
  if (exp == 0) 1 else base * pow(base, exp)
```

`@inline` denotes that the `pow` function it self must be inlined at application and `@ct` requires that the `exp` argument is a compile-time view of `Int`. The application of the function `pow` with a constant exponent will produce:

```
pow(base, 4)
  ↪ base * base * base * base * 1
```

Here, in the function application, constant 4 is promoted to `ct` by the automatic conversions. **[TODO: ref]**

3.3 Variable Argument Functions

Variable argument functions appear in widely used languages like Java, C#, and Scala. Such arguments are typically passed in the function body inside of the data structure (*e.g.* `Seq[T]` in Scala). When applied with variable arguments the size of the data-structure is statically known and all operations on them can be partially evaluated. However, sometimes, the function is called with arguments of dynamic size. For example, function `min` that accepts multiple integers

```
def min(vs: Int*): Int = vs.tail.foldLeft(vs.head) {
  (min, el) => if (el < min) el else min
}
```

Table 1. Types and corresponding method signatures after the translation to the compile-time view.

Annotated Type	Type's Method Signatures
`Int@ct`	`+(rhs: Int@ct): Int@ct`
`Vector[Int]@ct`	`map[U](f: (Int => U)@ct): Vector[U]@ct`
	`length: Int@ct`
`Vector[Int@ct]@ct`	`map[U](f: (Int@ct => U)@ct): Vector[U]@ct`
	`length: Int@ct`
`Map[Int@ct, Int]@ct`	`get(key: Int@ct): Option[Int]@ct`

Table 2. Promotion of terms to their compile-time views.

Promoted Term	Term's Promoted Type
`ct(Vector)(1, 2, 3)`	`: Vector[Int]@ct`
`ct(Vector)(ct(1), ct(2), ct(3))`	`: Vector[Int@ct]@ct`
`new (Cons@ct)(1, Nil)`	`: Cons[Int]@ct`
`new (Cons@ct)(ct(1), ct(Nil))`	`: Cons[Int@ct]@ct`
`ct((x: Int) => x)`	`: (Int@ct => Int@ct)@ct`
`inline((x: Int) => x)`	`: (Int => Int)@ct`

```
def min(vs: Int*): Int = macro
  if (isVarargs(vs)) q"min_CT(vs)"
  else q"min_D(vs)"

def min_CT(vs: Seq[Int] @ct): Int =
  vs.tail.foldLeft(vs.head) { (min, el) =>
  if (el < min) el else min
}
def min_D(vs: Seq[Int]): Int =
  vs.tail.foldLeft(vs.head) {
    (min, el) => if (el < min) el else min
  }
```

Figure 2. Function `min` is desugared into a `min` macro that based on the binding time of the arguments dispatches to the partially evaluated version (`min_CT`) for statically known varargs or to the original min function for dynamic arguments `min_D`.

can be called either with statically known arguments (*e.g.* `min(1,2)`) or with dynamic arguments:

```
val values: Seq[Int] = ... // dynamic value
min(values: _*)
```

Ideally, we would be able to achieve partial evaluation if the arguments are of statically known size and avoid partial evaluation in case of dynamic arguments. To this end we translate the method `min` into a partially evaluated version and a dynamic version. The call to these methods is dispatched, at compile-time, by the `min` method which checks if arguments are statically known. Desugaring of `min` is shown in Figure 2.

3.4 Removing Abstraction Overhead of Type-Classes

[TODO: not-sure how to achieve this!] [TODO: cite] Type-classes are omnipresent in everyday programming as they provide allow abstraction over generic parameters

```
object Numeric {
  @inline implicit def dnum: Numeric[Double] =
    DoubleNumeric
  @inline def zero[T](implicit num: Numeric[T]): T =
    num.zero
}

trait Numeric[T] {
  def plus(x: T, y: T): T
  def times(x: T, y: T): T
  def zero: T
}

class DoubleNumeric[T <: Double] extends Numeric[Double] {
  @inline def plus(x: T, y: T): T = x + y
  @inline def times(x: T, y: T): T = x * y
  @inline def zero: T = 0.0
}
```

Figure 3. Function for computing the non-negative power of a real number.

(*e.g.* Numeric abstracts over numeric values). Unfortunately, type-classes are a source of abstraction overheads during execution **[TODO: cite]**. Type-classes are in most of the cases statically known. Ideally, we would be able to deterministically remove abstraction overheads of type classes.

3.5 Dot Product

- Explain the removal of type classes together with inline. Explain how type classes are @i? and how they will completely evaluate if they are passed a static value.

- Comparison to other approaches.

4. Limitations

[TODO: discuss]

- Interaction with type variables.

- Type variables.
- Type annotations and overloading and implicit search.

5. The $F_{i<:}$ Calculus

We formalize the essence of our inlining system in a minimalistic calculus based on $F_{<:}$ with lazy records. To accommodate predictable partial evaluation we introduce binding-time annotations into the type system as first-class types that represent three kinds of bindings:

1. **Dynamic binding**. These are the types which express computation at runtime. All types written in the end user code are considered to be dynamic by default if no other binding-time annotation is given.

2. **Static binding**. Values of static terms can be computed at compile-time (*e.g.*constant expressions) but at are still evaluated at runtime by default. All language literals are static by default.

3. **Inline binding**. And finally the types that correspond to terms that are hinted to be computed at compile-time whenever possible.

5.1 Composition

An interesting consequence of encoding of binding times as first-class types is ability to represent values which are partially static and partially dynamic.

For example lets have a look at simple record that describes a complex number with two possible representations encoded through $isPolar$ flag:

$$complex : static\, \{isPolar : static\, Boolean,\, a : Double,\, b : Double\} \in \Gamma$$

This type is constructed out of a number of components with varying binding times. Representation encoding is known in advance and is static according to the signature. Coordinates a and b do not have any binding-time annotation meaning that they are dynamic.

Given this binding to $complex$ in our environment Γ we can use $inline$ to obtain a compile-time view to evaluate acess to $isPolar$ field at compile-time:

$$inline\, complex.isPolar : inline\, Boolean$$

Any statically known expression can be promoted via $inline$. Selection of dynamic fields on the other hand will return dynamic values despite the fact that record is statically known. In practice this can be used to specialize a particular execution path in the application to a particular representation by selectively inlining statically known parts.

Once you have inline view of the term it's also possible to demote it back to runtime evaluation through $dynamic$ view.

Not all type and binding time combinations are correct though. We restrict types to disallow nesting of more specific binding times into less specific ones.

$$\text{wff } iAny \qquad \text{(W-ANY)}$$

$$\frac{i <: j \quad i <: k \quad \text{wff } jT_1 \quad \text{wff } kT_2}{\text{wff } i(jT_1 \Rightarrow kT_2)} \qquad \text{(W-ABS)}$$

$$\frac{i <: j \quad i <: k \quad \text{wff } jS \quad \text{wff } kT}{\text{wff } i([X <: jS] \Rightarrow kT)} \qquad \text{(W-TABS)}$$

$$\frac{\forall j. \quad i <: j \quad \overline{\text{wff } jT}}{\text{wff } i\{\overline{x : jT}\}} \qquad \text{(W-REC)}$$

Figure 5. Well formed types wff iT

This restiction allows us to reject programs that have inconsistent annotations. For example the following function has incorrectly annotated parameter binding time:

$$(x : inline\, Int) \Rightarrow x + 1$$

This is inconsistent because the body of the function might not be evaluated at compile-time (as the function is not inline.) As described in (W-ABS) functions may only have parameters that are at most as specific as the function binding-time. In our example this doesn't hold as $inline$ is more specific than implicit $static$ annotation on function literal.

5.2 Subtyping

Another notable feature of our binding-time analysis system is deep integration with subtyping. We believe that such integration is crucial for an object-oriented language that wants to incorporate partial evaluation.

At core of the subtyping relation we have a subtyping relation on binding-time information with $dynamic$ as top binding-time.

$$i <: dynamic \qquad\qquad inline <: static$$
$$\text{(I-DYNAMIC)} \qquad\qquad \text{(I-STATIC2)}$$
$$static <: static \qquad\qquad inline <: inline$$
$$\text{(I-STATIC1)} \qquad\qquad \text{(I-INLINE)}$$

Figure 6. Binding-time subtyping.

We proceed by threading binding time information throughout regular $F_{<:}$ subtyping rules augmented with standard record types.

$t ::=$	Terms:
x, y	identifier
$(x : iT) \Rightarrow t$	function
$t(t)$	application
$\{\overline{x = t}\}$	record
$t.x$	selection
$[X <: iT] \Rightarrow t$	type abstraction
$t[iT]$	type application
$inline\ t$	inline view
$v ::=$	Values:
$x \Rightarrow t$	function value
$\{\overline{x = t}\}$	record value

$S, T, U ::=$	Types:
$iS \Rightarrow jT$	function type
$\{\overline{x : iS}\}$	record type
$[X <: iS] \Rightarrow jT$	universal type
Any	top type
$iT, jT, kT, lT ::=$	Binding-Time Types:
X	type identifier
$T, dynamic\ T$	dynamic type
$static\ T$	static type
$inline\ T$	inline type
$\Gamma ::=$	Contexts:
$\emptyset$	empty context
$\Gamma,\ x : iT$	term binding
$\Gamma,\ X <: iT$	type binding

Figure 4. Syntax of $F_{i<:}$

$$\Gamma \vdash iS <: Any \qquad \text{(S-Top)}$$

$$\Gamma \vdash iS <: iS \qquad \text{(S-Refl)}$$

$$\frac{\Gamma \vdash iS <: jU \quad \Gamma \vdash jU <: kT}{\Gamma \vdash iS <: jU} \quad \text{(S-Trans)}$$

$$\frac{i <: j \quad \Gamma \vdash S <: T}{\Gamma \vdash iS <: jT} \quad \text{(S-Inline)}$$

$$\frac{X <: iT \in \Gamma}{\Gamma \vdash X <: iT} \quad \text{(S-TVar)}$$

$$\{x_p : i_p T_p {}^{p \in 1..n+m}\} <: \{x_p : i_p T_p {}^{p \in 1..n}\} \quad \text{(S-Width)}$$

$$\frac{\Gamma \vdash kT_1 <: iS_1 \quad \Gamma \vdash jS_2 <: lT_2}{\Gamma \vdash iS_1 \Rightarrow jS_2 <: kT_1 \Rightarrow lT_2} \quad \text{(S-Arrow)}$$

$$\frac{\forall p \in 1..n.\ i_p S_p <: j_p T_p}{\{x_p : i_p S_p {}^{p \in 1..n}\} <: \{x_p : j_p T_p {}^{p \in 1..n}\}} \quad \text{(S-Depth)}$$

$$\frac{\Gamma, X <: iU_1 \vdash jS_2 <: kT_2}{\Gamma \vdash [X <: iU_1] \Rightarrow jS_2 <: [X <: iU_1] \Rightarrow kT_2} \quad \text{(S-All)}$$

$$\frac{\{x_p : i_p S_p {}^{p \in 1..n}\}\ \text{is permutation of}\ \{y_p : j_p T_p {}^{p \in 1..n}\}}{\{x_p : i_p S_p {}^{p \in 1..n}\} <: \{y_p : j_p T_p {}^{p \in 1..n}\}}$$
$$\text{(S-Perm)}$$

Figure 7. Subtyping.

Integration between binding-time subtyping and subtyping on regular types is expressed through (S-Inline) rule that merges the two into one coherent relation on binding-time types.

5.3 Generics

Crucial consequence of our design choices made in the system manifests in ability to use regular generics as means to abstract over binding-time without any additional language constructs.

For example given a generic identity function:

$$identity : static\ ([X <: Any] \Rightarrow static\ (X \Rightarrow X)) \in \Gamma$$

We can instantiate it to both in static and dynamic contexts through corresponding type application:

$$identity[static\ Int] : static\ (static\ Int \Rightarrow static\ Int)$$
$$identity[Int] : static\ (Int \Rightarrow Int) \qquad (1)$$

In practice this allows us to write code that is polymorphic in the binding time without any code duplication which is quite common in other partial evaluation systems.

This is possible due to the fact that we've integrated binding time information into types and augmented subtyping relation with subtyping

5.4 Typing

To enforce well-formedness of types in a context of partial evaluation we customize standard typing rules with additional constraints with respect to binding time.

$$\frac{x : iT \in \Gamma}{\Gamma \vdash x : iT} \quad \text{(T-IDENT)}$$

$$\frac{\forall t. \quad \Gamma \vdash t : jT \quad \text{wff } i\{x : jT\}}{\Gamma \vdash i\{\overline{x = t}\} : i\{\overline{x : jT}\}} \quad \text{(T-REC)}$$

$$\frac{\Gamma \vdash t_1 : i(jT_1 \Rightarrow kT_2) \quad \Gamma \vdash t_2 : jT_1}{\Gamma \vdash t_1(t_2) : kT_2} \quad \text{(T-APP)}$$

$$\frac{\Gamma \vdash t : i\{x = jT_1, \overline{y = kT_2}\}}{\Gamma \vdash t.x : jT_1} \quad \text{(T-SEL)}$$

$$\frac{t \text{ is not literal} \quad \Gamma \vdash t : static\, T}{\Gamma \vdash inline\, t : inline\, T} \quad \text{(T-INLINE)}$$

$$\frac{t \text{ is not literal} \quad \Gamma \vdash t : iT}{\Gamma \vdash dynamic\, t : dynamic\, T} \quad \text{(T-DYNAMIC)}$$

$$\frac{\Gamma \vdash t : iS \quad \Gamma \vdash iS <: jT}{\Gamma \vdash t : jT} \quad \text{(T-SUB)}$$

$$\frac{\Gamma, x : jT_1 \vdash t : kT_2 \quad \text{wff } i(jT_1 \Rightarrow kT_2)}{\Gamma \vdash i((x : jT_1) \Rightarrow t) : i(jT_1 \Rightarrow kT_2)} \quad \text{(T-FUNC)}$$

$$\frac{\Gamma, X <: jT_1 \vdash t_2 : kT_2 \quad \text{wff } i([X <: jT_1] \Rightarrow kT_2)}{\Gamma \vdash i([X <: jT_1] \Rightarrow t_2) : i([X <: jT_1] \Rightarrow kT_2)}$$
$$\text{(T-TABS)}$$

$$\frac{\Gamma \vdash t_1 : i([X <: jT_{11}] \Rightarrow kT_{12}) \quad \Gamma \vdash lT_2 <: jT_{11} \quad \Gamma \vdash i <: l}{\Gamma \vdash t_1[lT_2] : [X \mapsto lT_2]kT_{12}}$$
$$\text{(T-TAPP)}$$

Figure 8. Typing.

The most significant changes lie in:

- Additional checks in literal typing that ensure that constructed values correspond to well-formed types (T-FUNC, T-REC, T-TABS). To do this we typecheck literals together with possible binding-time term that might enclose it.

- New typing rules for binding-time views (T-INLINE, T-DYNAMIC). These rules only cover non-literal terms as composition of binding-time view and literal itself is handled in corresponding typing rule for given literal.

5.5 Partial Evaluation

In order to simplify partial evaluation rules we erase all of the type information before partial evaluation. This means that all functions become function values, type abstraction and application are complete eliminated.

Figure 11. Compilation pipeline.

$$\frac{t \rightsquigarrow t'}{x \Rightarrow t \rightsquigarrow x \Rightarrow t'} \quad \text{(PE-FUNC)}$$

$$\frac{\overline{t} \rightsquigarrow \overline{t'}}{\{\overline{x = t}\} \rightsquigarrow \{\overline{x = t'}\}} \quad \text{(PE-REC)}$$

$$\frac{t_1 \rightsquigarrow t'_1 \quad t'_1 \neq inline\, x \Rightarrow t \quad t_2 \rightsquigarrow t'_2}{t_1(t_2) \rightsquigarrow t'_1(t'_2)} \quad \text{(PE-APP)}$$

$$\frac{t_1 \rightsquigarrow inline\, x \Rightarrow t \quad t_2 \rightsquigarrow t'_2 \quad [x \mapsto t'_2]t \rightsquigarrow t'}{t_1(t_2) \rightsquigarrow t'} \quad \text{(PE-IAPP)}$$

$$\frac{t \rightsquigarrow t' \quad t' \neq inline\, x \Rightarrow t}{t.x \rightsquigarrow t'.x} \quad \text{(PE-SEL)}$$

$$\frac{t \rightsquigarrow inline\, \{x = t_x, \overline{y = t_y}\} \quad t_x \rightsquigarrow t'_x}{t.x \rightsquigarrow t'_x} \quad \text{(PE-ISEL)}$$

$$\frac{t \text{ is not literal} \quad t \rightsquigarrow t' \quad t' \Downarrow v}{inline\, t \rightsquigarrow inline\, v} \quad \text{(PE-INLINE)}$$

Figure 9. Partial evaluation $t \rightsquigarrow t'$

5.6 Evaluation

Once partial evaluation is complete we strip all binding-time terms and use regular untyped lambda calculus evaluation rules extended with lazy records.

$$v \Downarrow v \quad \text{(E-VALUE)}$$

$$\frac{t_1 \Downarrow x \Rightarrow t \quad t_2 \Downarrow v \quad [x \mapsto v]t \Downarrow v'}{t_1(t_2) \Downarrow v'} \quad \text{(E-APP)}$$

$$\frac{t \Downarrow \{x = t_x, \overline{y = t_y}\} \quad t_x \Downarrow v}{t.x \Downarrow v} \quad \text{(E-SEL)}$$

Figure 10. Evaluation $t \Downarrow v$

5.7 Conjectures

1. Progress.

2. Preservation.

3. Static terms are closed over statically bound variables.

4. Inline terms will be replaced with canonical value of corresponding type after partial evaluation.

6. Integrating $F_{i<:}$ with Object Oriented Languages

The $F_{i<:}$ calculus §5 captures the essence of user-controlled predictable partial-evaluation. In practice, though, it is fairly low level and it is not obvious how to define *classes* and methods from in modern multi-paradigm programming languages. Furthermore, $F_{i<:}$ requires an inconveniently large

$$\frac{\Pi \vdash T \in \Pi}{\Pi \vdash iT \rightsquigarrow iT}$$
$$\text{(CT-TVAR)}$$

$$\frac{\Pi \vdash T \notin \Pi}{\Pi \vdash iT \rightsquigarrow inline\ T}$$
$$\text{(CT-T-VAR)}$$

$$\frac{\overline{\Pi \vdash t \rightsquigarrow t'}}{\Pi \vdash i\overline{\{x = t\}} \rightsquigarrow inline\ \overline{\{x = t'\}}} \quad \text{(CT-REC)}$$

$$\frac{\overline{\Pi \vdash iT \rightsquigarrow jT}}{\Pi \vdash \overline{\{x : iT\}} \rightsquigarrow inline\ \overline{\{x : jT\}}} \quad \text{(CT-T-REC)}$$

$$\frac{\Pi \vdash iT \rightsquigarrow jT \quad \Pi \vdash kS \rightsquigarrow lS}{\Pi \vdash iT \Rightarrow kS \rightsquigarrow jT \Rightarrow lS} \quad \text{(CT-T-ARROW)}$$

$$\frac{\Pi \vdash jT \rightsquigarrow kT}{\Pi \vdash [X <: iS] \Rightarrow jT \rightsquigarrow [X <: iS] \Rightarrow kT} \quad \text{(CT-T-UNIV)}$$

$$\frac{\Pi \vdash t \rightsquigarrow t' \quad \Pi \vdash iT \rightsquigarrow jT}{\Pi \vdash i(x : iT) \Rightarrow t \rightsquigarrow inline\ (x : jT) \Rightarrow t'} \quad \text{(CT-FUNC)}$$

$$\frac{\Pi, X \vdash t \rightsquigarrow t'}{\Pi \vdash i([X <: jT_1] \Rightarrow t) \rightsquigarrow inline\ ([X <: jT_1] \Rightarrow t')}$$
$$\text{(CT-TABS)}$$

$$\frac{\Pi \vdash t \rightsquigarrow t' \quad \Pi \vdash iT \rightsquigarrow jT}{\Pi \vdash t[iT] \rightsquigarrow t'[jT]} \quad \text{(CT-TAPP)}$$

Figure 13. Translation of a type abstractions, function, and record values into a compile-time view. The translation is used for promoting types into their compile time versions.

number of `inline` calls in method invocations. In this section we a scheme for translating classes into $F_{i<:}$ (§6.1), show how to provide compile time views of classes and and *methods*§**??**, and formalize convenient implicit conversions for the calculus §6.3.

Furthermore, rules of $F_{i<:}$ do not support effect-full computations and each `inline` term is trivially converted to a dynamic term after erasure. In case of languages that do support mutable state and side-effects this needs to be treated specially. For simplicity, we omit side-effects from our discussion and assume that all partially evaluated code is side-effect free and that each `inline` term can be converted to dynamic code.

6.1 Desugaring Object Oriented Constructs to $F_{i<:}$

6.2 Compile-Time View of the Terms

6.3 Implicit Conversions

According to $F_{i<:}$ rules if method signatures contain compile-time views of a type the corresponding arguments in method application would always have to be promoted to `inline`. In practice this is not convenient as it requires an inconveniently large number of annotations. Partial evaluation is an optimization, and as such, it should not affect user code - users should not be aware of the internal operation of the library.

To address this issue we introduce implicit conversions from all language literals, and direct class constructor calls of non-inline type into their compile-time views. For example, for a factorial function

```
def fact(n: Int @ct) = if (n == 0) 1 else fact(n - 1)
```

we will not require annotations on literals 0, and 1. Furthermore, the function can be invoked without promoting the literal 5 into it's compile-time view:

```
fact(5)
  ↪ 120
```

7. Evaluation

7.1 Reduction of Code Duplication

7.2 Performance Comparison

Table 3. Performance comparison with LMS and hand optimized code.

Benchmark	Hand Optimized	LMS	Scala Inline
`pow`			
`min`			
`dot`			
`fft`			

8. Related Work

9. Conclusion

References

[1] Jacques Carette and Oleg Kiselyov. Multi-stage programming with functors and monads: Eliminating abstraction overhead from generic code. In *Generative Programming and Component Engineering (GPCE)*, 2005.

[2] Olivier Danvy. *Type-directed partial evaluation*. Springer, 1999.

[3] Yoshihiko Futamura. Partial evaluation of computation process–an approach to a compiler-compiler. *Higher-Order and Symbolic Computation*, 12(4):381–391, 1999.

[4] Neil D. Jones, Carsten K. Gomard, and Peter Sestoft. *Partial Evaluation and Automatic Program Generation*. Prentice Hall, 1993.

[5] Manohar Jonnalagedda, Thierry Coppey, Sandro Stucki, Tiark Rompf, and Martin Odersky. Staged parser combinators for efficient data processing. In *International Conference on Object Oriented Programming Systems Languages and Applications (OOPSLA)*, 2014.

[6] V. Jovanovic, A. Shaikhha, S. Stucki, V. Nikolaev, Koch C., and M Odersky. Yin-yang: Concealing the deep embedding of DSLs. In *International Conference on Generative Programming and Component Engineering (GPCE)*, 2014.

[7] Stefan Monnier and Zhong Shao. Inlining as staged computation. *Journal of Functional Programming*, 13(03):647–676, 2003.

[8] Martin Odersky and Konstantin Läufer. Putting type annotations to work. In *Symposium on Principles of Programming Languages (POPL)*, 1996.

$$\llbracket \textbf{let } x : T_x = t_x \textbf{ in } t \rrbracket = ((x : T_x) \Rightarrow t)(t_x)$$

$$\llbracket \textbf{let } type\ T_1 = T_2 \textbf{ in } t \rrbracket = ([T_1 <: T_2] \Rightarrow t)[T_2]$$

$$\llbracket \textbf{let } class\ C[A](x : T_x)\{def\ f[B](y : T_y) = t_f\} \textbf{ in } t \rrbracket =$$

$$\textbf{let } type\ C = [A] \Rightarrow inline\ \{fields : \{x : T_x\}, methods : inline\ \{f : [B] \Rightarrow T_y \Rightarrow T_f\}\} \textbf{ in}$$

$$\textbf{let } C : [A] \Rightarrow inline\ ((t_x : T_x) \Rightarrow C[A]) = [A] \Rightarrow inline\ ((t_x : T_x) \Rightarrow$$

$$inline\ \{fields = \{x = t_x\}, methods = inline\ \{f = [B] \Rightarrow (y : T_y) \Rightarrow t_f\}\}) \textbf{ in } t$$

Figure 12. Desugaring of classes into $F_{i<:}$.

[9] Bruno CdS Oliveira, Adriaan Moors, and Martin Odersky. Type classes as objects and implicits. In *ACM Sigplan Notices*, volume 45, pages 341–360, 2010.

[10] Tiark Rompf and Martin Odersky. Lightweight modular staging: a pragmatic approach to runtime code generation and compiled DSLs. *Communications of the ACM*, 55(6):121–130, June 2012.

[11] Tiark Rompf, Arvind K. Sujeeth, Kevin J. Brown, HyoukJoong Lee, Hassan Chafi, Kunle Olukotun, and Martin Odersky. Project Lancet: Surgical precision JIT compilers. In *International Conference on Programming Language Design and Implementation (PLDI)*, 2013.

[12] Arvind K Sujeeth, Austin Gibbons, Kevin J Brown, HyoukJoong Lee, Tiark Rompf, Martin Odersky, and Kunle Olukotun. Forge: generating a high performance DSL implementation from a declarative specification. In *International Conference on Generative Programming and Component Engineering (GPCE)*, 2013.

[13] Walid Taha and Tim Sheard. Multi-stage programming with explicit annotations. In *Workshop on Partial Evaluation and Program Manipulation (PEPM)*, 1997.